

# Requirements

- ASM Bytecode Instrumentation Library 3.3:

- <http://www.cs.tufts.edu/research/redline/elephantTracks/asm-all-3.3.jar>

- IBM's thread building blocks (including header files):

- <http://threadingbuildingblocks.org/download#stable-releases>

## Building Elephant Tracks

1. First, copy Makefile.inc to Makefile.local; in Makefile.local you will set local build parameters for your machine, in particular, you must set the following

- INSTALL\_DIR -- Where you would like the Elephant Tracks binaries be installed
- ASMJAR -- Location of the asm.jar ( asm version > 3.3)
- JAVA\_PATH -- Must point to the home directory for java, not the actually java binary. (ie, the 'java' binary should be in JAVA\_PATH/bin/java)

2. make; make install

### Known Issue

There is disagreement between different versions of jni.h on the const-ness of certain parameters. This can lead to the following error when attempting to build Elephant Tracks:

**ETCallbackHandler.cpp: In function '\_jobject\* instNewObjectA(JNIEnv\*, jclass, jmethodID, const jvalue\*)':**  
**ETCallbackHandler.cpp:1385:66: warning: invalid conversion from 'const jvalue\*' to 'jvalue\*'**

This can be resolved by removing the -Werror g++ flag, and adding -fpermissive.

# Running Elephant Tracks

First, make sure that `INSTALL_DIR` is in the `LD_LIBRARY_PATH`

Then run any Java program with Elephant Tracks like so:

```
java -classpath <other-paths>:$ASMJAR -Xbootclasspath/a:$INSTALL_DIR \
-agentlib:ElephantTracks=<ElephantTracks Options>
```

## Options:

Are represented as name=value pairs, and may be given on the command like so:

```
-agentlib:ElephantTracks==name1=value1@name2=value2@...@nameN=valueN
```

Note that JVMs often impose an (undocumented) limit on the length of the command strings passed to a JVTI agent such as Elephant Tracks, and will silently truncate if this limit is exceeded. For this reason, it is recommended that infrequently changing options are stored in an option file, and specify that option file on the command line (see below).

- `optionsFile=<option file>`

- This is a path to a file containing options (the same as may be passed on the command line), one option per line.

- `classPath=<path>`

This is the path ElephantTracks will use to start its own Java process (not the one running your program). It must include `INSTALL_DIR`, `INSTALL_DIR/elephantTracksRewriter.jar`, and the `asm-3.3 jar` file.

- `javaPath=<path to java executable>`

- This is the path to the actual java binary, not merely the directory it is in. (for example, `/usr/bin/java`, not `/usr/bin`).

- `namesFile=<file name>`

- The file in which to output the names information (see below)

- `traceFile=<file name>`

- The file in which to output the trace. You may also redirect trace output to a shell command with this syntax:

traceFile=>(shell command)

For example:

traceFile=>(gzip > foo.trace.gz)

Would pipe the trace output to gzip, and redirect the output from gzip to foo.trace.gz

- bufferSize=<number>

- How many records to hold in Elephant Tracks' internal buffer; larger values

generally give better performance, but will use more memory (approximately 40 bytes per record).

## Record Types

### Names File Records

The entries in the names file map names of methods to numerical ids used in the trace.

### # Comment

Lines that start with # are comments.

### C 0xcccc class-name [0xssss] (I:0xiiii)\*

cccc = class id; ssss = superclass id;

iiii = superinterface id

### I 0xiiii interface-name [0xssss] (I:0xiiii)\*

C or I will be output when the id is assigned, which is on first \*mention\*, and will be output again, possibly with more info, such as superclass / superinterfaces, when \*processed\*. We could perhaps distinguish a mention from a definition if you think it would be easier ...

### E 0xcccc class-name

marks end of processing a class/interface \*definition\*

(so now you know all the methods and fields)

F I/S 0xffff name 0xcxxx class-name field-type

I = instance; S = static

ffff = field id; cxxx = declaring class id

field-type is a descriptor (I, etc., or Jclassname; etc.)

## **N 0xmmmm 0xcxxx class-name method-name descriptor flags**

mmmm = method id; cxxx = class id

flags can include I or S for instance or static,  
with N added if native

## **S 0xmmmm 0xcxxx 0xssss descriptor dims**

For allocation sites

mmmm = method id; cxxx = declaring class id;

ssss = site id (unique in run)

descriptor = type of thing allocated

dims = number of dimensions (0 for a scalar object)

## **Trace File Records**

The entries in the trace file represent events in the program execution.

### **# Comment**

Lines that start with # are comments.

### **Allocation:**

There are several different kinds of allocation record, based on how the object was allocated. They all have the same general form:

**<Type Character> <object-id> <size> <type>  
<site> <length> <thread-id>**

### **Types:**

**A** - Array allocation

**I** - Initial Heap Allocation (Allocated before Elephant Tracks started)

**N** - Allocated via the NEW byte code in the traced java program

**P** - Preexisting object; these are objects for which Elephant Tracks missed the actual object allocation, but discovered later. This can be due to objects from the constant pool, or VM bugs.

**V** - Objects allocated by the virtual machine.

The new object has ID object-id, which is used to refer to the object in later events; the size in bytes; the type (a Java type as a string); length field is 0 for non-arrays, and the length of the array for arrays, the ID of the allocating thread.

The size is the size of the object in bytes as reported by the VM (this includes object headers and possibly other VM structures).

## **Death:**

**D <object-id> <thread-id>**

Object-id died, and its death occurred in thread\_id.

## **Field update:**

**U <old-target-id> <object-id>**

**<new-target-id> <field-id> <thread-id>**

The field field-id in object object-id that used to point to old-target-id now points at new-target-id, and this update occurred in thread-id.

An object-id of 0 indicates that this is an update to a static field.

## **Method entry:**

**M <method-id> <receiver-object-id>**

## **<thread-id>**

A call to the method method-id with receiver object receiver-object-id in thread thread-id.

A origin receiver of 0 indicates that this was a static method

## **Method exit:**

**E <method-id> <receiver-object-id>  
<thread-id>**

Return from method method-id with receiver object receiver-object-id in thread thread-id.

A receiver of 0 indicates this was a static method.

**X <method-id> <receiver-object-id>  
<exception-id> <thread-id>**

Exceptional exit from a method method-id, with receiver object receiver-object-id, exception object exception-id, in thread thread-id

## **Exception Throw**

**T <method-id> <receiver-id>  
<exception-object-id> <thread-id>**

An exception was thrown in the given method, with the given receiver, and with the exception object itself having the given id.

## **Exception Handled**

**H <method-id> <receiver-id>**

## **<exception-object-id> <thread-id>**

An exception was handled in the given method, with the given receiver, and with the exception object itself having the given id.

## **R <root object ID> <thread id>**

Object-ID was a root in thread-id